


Dollhouse


Progress Update #1

Leonhard Driesch

Visual Computing & Artificial Intelligence
Prof. Matthias Nießner
Technical University of Munich

September 23rd, 2026





Gemma 4 12B. Finds and boxes objects in an image. Change the prompt to look for specific things, e.g. detect chairs and lamps.


Input image







Detections






Prompt

detect all objects

Detect

Boxes (pixels)

label	x1	y1	x2	y2
sofa	877	493	1272	811
sofa	295	507	631	852
cabinet	1229	449	1575	757
table	672	571	897	800

Detection: which model?

Find every object in a room photo, no fixed class list

VLM

current attempt

Gemma 4 12B, prompt “detect all objects”, boxes parsed from the text output.

- + Free-form prompt; can decide what counts as an object
- Boxes come from generated text
- Large (12B) and slow

Open-vocabulary detector

Grounding DINO, OWLv2, YOLO-World, SAM 3, ...

- + Trained for localization; masks with SAM 3
- + Small and fast
- Needs the class names as text input

VLM, detector, or both (VLM names the objects, detector localizes them)?

Other models for this use case?

Prompt

Extract the couch from this image as a standalone picture with a transparent background.

Input images (optional)



Aspect ratio

Auto uses 1:1 when there is no input image.

Auto (from first input image) ▾

Size

☐ Full (~4 MP)
 ☒ Half (~1 MP, faster)

Steps

10

40


60

Seed

☒ Random seed

41782881

Result



Send to 3D tab →

Input image







Resolution

Voxel grid size. 512 is fastest; 1536 is most detailed.

☐ 512
 ☒ 1024
 ☐ 1536

Target face count





Texture size

☐ 1024
 ☒ 2048
 ☐ 4096

Seed

☒ Random seed

Generate 3D model

Result



Subject as seen by the model



Download GLB

20260923-... .glb

15.5 MB ↓



Original photo



Blender render by GPT-6 Astra

Next steps

- Literature review
- Get detection working
- Scene assembly

Goal: an end-to-end baseline (MVP)